



GLOBAL GOVERNANCE FRAMEWORKS · TIER 3 CROSS-CUTTING
PROTOCOL

Adaptive Coordination Protocol

*A lightweight, voluntary coordination mechanism designed to align
national sovereignty with global systemic needs during cross-border crises.*

V0.5

By **Björn Kenneth Holmström**

with Claude, ChatGPT and DeepSeek

Status: Ready for Practitioner Review

Adaptive Coordination Protocol (ACP)

Tier 3 (Cross-Cutting Protocol)

Status: Draft v0.5 – Refinement & Realism Upgrade

Purpose: Real-World Coordination Under Systemic Stress

0. Revision Note (v0.5)

This revision addresses targeted refinements identified in v0.4 review:

- **Narrative fragmentation metric** complemented by direct community trust surveys
 - **Constrained coordination mode** concretely defined
 - **Rapid Narrative Assembly fallback** specified for crisis conditions
 - **Domain Adaptation Protocol** strengthened with external stakeholder involvement
 - **Complexity acknowledgment** with commitment to future simplification review
-

1. Honest Framing: The Coordination Spectrum

1.1 Acknowledging Soft Power

ACP is **voluntary** in the sense that no state is legally compelled to join or comply with **allocation proposals**. However, it employs **soft power mechanisms**—transparency,

reputation, access to benefits—that create incentives for participation and costs for defection.

These mechanisms exist on a spectrum:

Mechanism	Description	Nature
Transparency	Public visibility of commitments and compliance.	Soft accountability. States may experience reputational pressure but retain full sovereignty over actions.
Tiered Benefits	Access to ERF, data, and coordination support scaled by participation level.	Conditional cooperation. Benefits are earned through engagement; non-participation means foregoing opportunities, not facing sanctions.
Exclusion from Benefits	Observing Partners receive reduced or no access to certain ACP resources.	Boundary-setting. The system protects its integrity by limiting free-riding, not by punishing.
Documentation of Disruption	Public logging of actions that actively undermine coordination.	Truth-telling. The system names behavior without imposing penalties, leaving consequences to political and diplomatic channels.
Referral to Multilateral Bodies	In extreme cases, patterns of disruption may be reported to existing institutions (e.g., UN).	Leveraging existing structures. ACP does not create new enforcement powers.

ACP does not resolve the tension between voluntary coordination and the need for system integrity. It **manages** this tension through graduated, transparent mechanisms that preserve sovereignty while protecting the commons.

2. ERF Governance (Detailed Specification)

2.1 Capitalization

2.1.1 Seed Funding (Phase 1)

Initial ERF capitalization comes from:

- **Voluntary contributions** from founding ACP participants, with a **non-binding guidance** scaled to GDP (e.g., 0.001% of GDP as a suggested minimum).
- **Multilateral seed grants** (e.g., UN CERF, World Bank trust funds).
- **Philanthropic bridge funding** restricted to the first 24 months.

2.1.2 Sustained Capitalization (Phase 2+)

- A **progressive contribution formula** is adopted by the Participant Assembly, based on capacity and historical crisis exposure.
- Contributions are **pledged annually** with public reporting.
- A **minimum viable fund threshold** is established (e.g., sufficient to cover 30 days of Emergency Mode operations). If the fund falls below this threshold, the Participant Assembly is convened to address the shortfall.

2.1.3 Shortfall Protocol (Updated)

If contributions are insufficient to meet Emergency Mode demands:

1. The ACP Secretariat notifies all Core and Cooperating Partners.
2. A **special pledging round** is initiated with a 7-day window.
3. If shortfall persists, the Allocation Engine adjusts proposals to reflect available resources, and the system operates in a **"constrained coordination" mode**.

Constrained Coordination Mode – Operational Definition:

When ERF resources fall below the minimum viable threshold, the following adjustments take effect:

- **First priority:** Direct life-saving assistance (emergency shelter, food, medical care) is protected. All other funding streams are subject to proportional reduction.
- **Second priority:** Frontline state stabilization funding is reduced by a predefined percentage (e.g., 25%), with the understanding that this may increase pressure on those states and potentially accelerate onward movement.
- **Deprioritized:** Long-term integration support, capacity-building programs, and non-essential administrative functions are suspended until the fund is replenished.
- **Transparency:** A public notice is issued on the ACP dashboard stating: *"ACP is operating in Constrained Coordination Mode. Available resources are insufficient to meet all commitments. The following categories are currently deprioritized: [list]."*

These adjustments are **pre-announced and standardized**, not determined ad hoc during crisis.

2.2 Disbursement Governance

2.2.1 Approval Body

The **ERF Disbursement Panel** consists of **7 rotating members**:

- 2 from frontline/transit states
- 2 from destination states
- 2 from contributing states (non-frontline, non-destination)
- 1 independent expert (appointed by the Participant Assembly)

Rotation occurs every 12 months, with staggered terms to ensure continuity.

2.2.2 Decision Rules

- **Automatic disbursements** (pre-authorized Emergency Mode envelope): No panel approval required; funds released within 48 hours.
- **Expedited disbursements** (> 30% of fund or > \$X million): Panel approval required within 72 hours, with a **simple majority** decision.
- **Standard disbursements:** Panel approval within 7 days, with **qualified majority (5/7)**.

2.2.3 Transparency and Appeal

- All disbursement decisions are published with rationale within 24 hours.
- Affected states may appeal a denial to the **Peer Mediation Panel**, which issues a non-binding recommendation within 7 days.

2.3 Secretariat Accountability and Oversight

Rationale: The ACP Secretariat holds significant operational authority. Mechanisms are required to detect and address potential capture, bad faith, or persistent underperformance by the Secretariat itself.

2.3.1 ACP Ombudsperson

- An independent **ACP Ombudsperson** is appointed by the Participant Assembly for a **single, non-renewable 4-year term**.
- **Mandate:** To receive and investigate confidential complaints regarding Secretariat conduct, including allegations of bias, data manipulation, favoritism, or procedural violations.
- **Access:** Any participant state, accredited civil society organization, or Community Reflection Circle may submit a complaint.
- **Reporting:** The Ombudsperson issues a **public annual report** summarizing investigations and recommendations. Urgent matters may be reported directly to the Participant Assembly at any time.

2.3.2 External Audit Trigger

- Any **three (3) Core Partners** acting jointly may trigger an **extraordinary external audit** of the Secretariat's operations, finances, or data integrity.
- The audit is conducted by an independent body (e.g., a multinational audit firm or UN Office of Internal Oversight Services) selected by the Peer Mediation Panel.
- Findings are presented to the Participant Assembly within 60 days and made public (with appropriate redactions for sensitive information).

2.3.3 No-Confidence Process

- If the Ombudsperson's report or an external audit identifies **systemic misconduct or gross negligence**, a motion of no-confidence in the Secretariat leadership may be introduced in the Participant Assembly.

- Passage requires a **two-thirds majority** of Core and Cooperating Partners.
 - Upon passage, the Participant Assembly shall initiate a **leadership transition process** within 30 days.
-

3. Pluralistic Narrative Infrastructure

3.1 Principle: Narrative as Co-Creation, Not Management

The v0.3 framing of "narrative support triggered upon activation" risked treating public narrative as a top-down communications function. v0.4 repositions narrative infrastructure as a **pluralistic, participatory process**.

3.2 ACP Narrative Commons

3.2.1 During Normal Mode

- The ACP dashboard includes an open **"Voices from the Corridor"** section, aggregating perspectives from frontline communities, civil society, and local media (curated for relevance, not editorial control).
- **Narrative Stewardship Guilds** (per Meaning Infrastructure Protocol) are invited to document and share stories of coordination and its challenges, independent of ACP Secretariat oversight.

3.2.2 Upon Stress/Emergency Mode Activation

- A **Rapid Narrative Assembly** is convened within 72 hours, comprising:
 - Representatives from affected communities (**selected via sortition where feasible**)
 - Local journalists and cultural practitioners
 - ACP Secretariat (observer only)

Fallback When Sortition Is Not Feasible:

In acute crisis conditions where sortition is impractical (e.g., active conflict, displaced populations in transit, lack of communication infrastructure), the following fallback applies:

1. The ACP Secretariat, in consultation with the **Peer Mediation Panel**, identifies **3–5 trusted civil society organizations or community-based networks** already operating in the affected area.
 2. These organizations are invited to **nominate representatives** to the Rapid Narrative Assembly, with a mandate to reflect the diversity of affected voices (including marginalized groups).
 3. The composition of the Assembly and the selection process are **transparently documented** on the ACP dashboard, including an explicit statement when sortition was not feasible and why.
- The Assembly produces a "**Living Narrative Brief**" —a plain-language, evolving document that explains the crisis and coordination efforts from multiple perspectives. This brief is **not a single official story** but a **container for diverse voices**.
 - The brief is updated weekly and translated into regional languages.

3.2.3 Community Reflection Circles

- These occur **during and after** crisis response, not only post-escalation.
- Findings are fed back into the Narrative Brief and into protocol refinement via the Participant Assembly.

4. Due Process for "Disruptive Actor" Categorization

4.1 Principle

Labeling a state or actor as "Disruptive" has significant diplomatic and operational consequences. The process must be fair, transparent, and subject to appeal.

4.2 Designation Process

4.2.1 Evidence Gathering

- The ACP Secretariat compiles a **documented pattern of behavior** (e.g., evidence of weaponizing migration flows, active obstruction of coordination, disinformation campaigns).
- Evidence must be **corroborated by at least two independent sources** (e.g., UN agencies, regional bodies, accredited NGOs).

4.2.2 Notification and Response Period

- The actor in question receives a **formal notification** with the evidence and a **30-day period** to respond or remedy the behavior.
- During this period, the actor is classified as **"Under Review."**

4.2.3 Designation Decision

- After 30 days, the **Peer Mediation Panel** reviews the evidence and the actor's response.
- Designation as **"Disruptive Actor"** requires a **consensus decision** of the Panel.
- The designation is **time-limited** (e.g., 6 months) and subject to quarterly review.

4.2.4 Appeal

- The designated actor may appeal to the **Participant Assembly**, which can overturn the designation by a **two-thirds majority**.

4.3 Transparency

- All designations, evidence summaries (with appropriate sensitivity), and responses are published on the ACP dashboard.
-

5. Domain Transferability

5.1 Acknowledgment

ACP emerged from migration-focused simulations. Its mechanisms are **not universally applicable without adaptation**. v0.5 explicitly acknowledges this.

5.2 Domain Adaptation Protocol (Updated)

Before ACP is deployed in a new domain (e.g., climate finance, pandemic response, AI disruption), a **Domain Adaptation Review** is conducted:

1. **Stakeholder Mapping:** Identify key actors, incentives, and existing coordination mechanisms in the new domain.
2. **Mechanism Audit:** Assess which ACP components transfer cleanly and which require modification.
3. **Pilot Design:** Develop a domain-specific pilot with adjusted thresholds, allocation factors, and narrative infrastructure.
4. **External Stakeholder Involvement:** For genuinely new domains where no Participant Assembly yet exists, the review process **SHALL** include **external domain experts and representatives of affected communities** as non-voting advisors. Their input is documented and considered in the adaptation design.
5. **Participant Assembly Approval:** Once a domain-specific Assembly is constituted, the adapted protocol is submitted for formal approval.

This ensures that domain adaptation is not a closed, self-referential process.

5.3 Example: Climate Finance vs. Migration

ACP Component	Migration Applicability	Climate Finance Applicability	Adaptation Needed
SAL	High (flows, capacities)	Medium (funding flows, project impacts)	Different data sources; slower tempo
Allocation Engine	High (redistribution of people)	Low (allocation of funds is different)	Replace with transparent funding allocation criteria
ERF	High (rapid frontline support)	High (rapid climate disaster response)	Similar mechanism, different triggers
Narrative Infrastructure	High	High	Similar principles, different framing

6. Measurable Success Indicators

6.1 Operationalized Metrics

Success Criterion	Indicator	Measurement Method	Baseline/Target
Rapid financial response	Time from Emergency Mode activation to first ERF disbursement	ERF transaction logs	Target: < 72 hours
Improved mid-level compliance	% of Cooperating Partners meeting ≥70% of commitments	Contribution Registry	Target: Increase from baseline (TBD in pilot)
Coherent public narrative (structural)	Reduction in narrative fragmentation score	Media monitoring (e.g., topic modeling across outlets)	Target: 20% reduction in fragmentation during crisis vs. pre-ACP baseline
Coherent public narrative (trust-based)	Community trust in coordination efforts	Annual trust survey of affected populations (stratified random sample; includes questions on perceived fairness, transparency, and whether coordination is understood)	Target: ≥60% reporting "coordination efforts are fair and understandable"
External actor containment	Number of documented disruptive actions by designated actors	Secretariat log; independent verification	Target: Decreasing trend over 12-month periods
Transparent trust tracking	Contribution Registry update frequency and accuracy	Audit of registry against actual flows	Target: 100% of commitments updated within 7 days of reporting period
Participant satisfaction	Annual survey of Participant	Structured questionnaire	Target: ≥70% reporting "ACP improves"

Success Criterion	Indicator	Measurement Method	Baseline/Target
	Assembly members		coordination outcomes"

Note on Narrative Metrics: The fragmentation score measures structural discourse coherence; the community trust survey measures lived perception. Both are necessary; neither is sufficient alone.

6.2 Qualitative Assessment

In addition to quantitative metrics, an annual **Qualitative Impact Review** is conducted, incorporating:

- Community Reflection Circle summaries
- Narrative Brief analysis
- Case studies of specific crisis responses

This review is published alongside quantitative metrics.

7. Complexity Acknowledgment and Simplification Commitment

7.1 Acknowledgment

As ACP has matured through iterative revision, its specification has grown in detail and complexity. This is a natural consequence of addressing real-world implementation challenges. However, complexity imposes costs:

- Implementation friction for smaller or lower-capacity states
- Accessibility barriers for non-specialist audiences

- Increased surface area for legalistic gaming of rules

7.2 Simplification Review

ACP commits to a **periodic simplification review**, conducted every **two years** by the Participant Assembly. The review asks:

What is the minimum viable version of each mechanism that achieves the core coordination goal?

The review may result in:

- Streamlining of procedural steps
- Consolidation of overlapping functions
- Removal of mechanisms that have proven unnecessary or unused

The outcome of each simplification review is published as a **Simplification Report**, and any resulting protocol changes follow the standard amendment process.

8. Updated Risk Profile

Risk	v0.5 Mitigation
Perception of coercion	Honest framing of soft power spectrum; transparency of all mechanisms
ERF undercapitalization	Shortfall protocol with clearly defined constrained coordination priorities
Narrative capture	Pluralistic Narrative Commons; sortition with documented fallback
Unfair "Disruptive" labeling	Due process with evidence, response period, peer review, and appeal
Domain mismatch	Domain Adaptation Protocol with external stakeholder involvement
Unmeasurable success	Operationalized indicators with baselines and targets, including community trust surveys
Complexity burden	Biennial simplification review commitment

9. Strategic Insight (v0.5)

This revision addresses the "last mile" of accountability and realism. The protocol now:

- **Defines what happens when resources are scarce** (constrained coordination mode)
- **Provides a credible fallback for narrative inclusion** when ideal sortition is impossible
- **Ensures domain adaptation is not an insular process** by involving external stakeholders
- **Commits to periodic simplification** to counter the natural drift toward complexity

The protocol is now ready for external practitioner review and pilot partner engagement.

10. Conclusion

ACP v0.5 is a **refinement release**. It does not introduce new structural components but strengthens the weak points identified in v0.4 review. The protocol is now:

- **Honest** about its power dynamics
- **Accountable** in its processes
- **Pluralistic** in its narrative infrastructure
- **Realistic** about resource constraints
- **Humble** about domain applicability
- **Committed** to managing its own complexity

Next Steps:

- Seek external review from migration coordination practitioners
- Identify a first-mover pilot partner (UN OCHA, EU, African Union)
- Develop a simplified "ACP Lite" one-pager for initial stakeholder conversations

Supplementary Addenda

Addendum A — Interpretation Containment Layer (ICL)

Addendum B — Pre-Action Commitment Layer (PACL)

Addendum C — System Regeneration Protocol (SRP)

Tier	3 (Cross-Cutting Protocol)
Supplements	ACP v0.5 — Refinement & Realism Upgrade
Status	Draft v0.1 — For practitioner review
Purpose	Address the pre-coordination gap: the failure modes that occur before and between formal ACP activation states

WHY *ACP v0.5 governs coordination under uncertainty. These addenda address an adjacent problem: the dynamics that determine whether coordination is even possible — how meaning forms under pressure, how action thresholds hold under stress, and how systems recover after they fracture.*

Addendum A — Interpretation Containment Layer (ICL)

Status	Draft v0.1 — Pre-Coordination Cognitive Stabilization
Supplements	ACP v0.5
Purpose	Slow and structure the formation of political meaning under high-stakes ambiguity, preventing premature narrative closure from triggering irreversible escalation before ACP coordination can engage

A.1 The Problem ACP v0.5 Does Not Yet Address

ACP v0.5 governs coordination once actors are operating within shared frameworks. It assumes that when a crisis event occurs, there is time for verification, deliberation, and allocation decisions to take place.

This assumption fails in a specific and important class of events: high-impact, attribution-ambiguous incidents — assassinations, precision strikes, covert operations — where the escalation timeline compresses to hours and narrative certainty forms faster than facts can settle.

In these conditions, escalation is not driven by events themselves but by the speed at which actors converge on a shared interpretation of who is responsible. Once that convergence occurs, retaliation becomes politically mandatory regardless of what the underlying facts are. ACP coordination becomes irrelevant before it can activate.

ICL addresses this gap directly.

A.2 Core Concept: The Legitimacy Delay Window

ICL is built around a single foundational mechanism: the Legitimacy Delay Window (LDW). The LDW is a predefined period following a high-impact incident during which:

- attribution claims remain formally probabilistic (not binary),
- retaliatory justification is classified as procedurally premature in ACP-linked channels,
- and escalation actions above a defined threshold cannot be legitimized by the system.

The LDW does not prevent speech, intelligence activity, or military readiness. It constrains only the formal closure of causal claims — the moment at which an actor officially declares certainty. This distinction is critical: ICL intervenes in the timing of meaning consolidation, not in the content of meaning itself.

A.3 Attribution Spectrum Engine (ASE)

All attribution claims processed through ACP-linked channels must be classified using a four-band spectrum:

Confidence Band	Threshold	Operational Meaning
Low	<40%	Multiple actors plausible; no dominant hypothesis
Medium	40–69%	Convergent suspicion; insufficient for action authorization
High (Strong Evidence)	70–84%	Substantial corroboration; eligible for limited diplomatic response
Verified	≥85%	Multi-source confirmation; full action legitimacy restored

Classification requires corroboration from at least two independent source categories (e.g., signals intelligence, human intelligence, third-party verification, physical evidence). Single-source claims are capped at Medium regardless of internal confidence levels.

A.4 Narrative Separation Layer

A key failure mode in acute crises is cascade certainty amplification: intelligence assessments are imported into political narratives, which are imported into media framing, which then feeds back into political demands for action — all within hours. Each transfer inflates confidence without adding evidence.

ICL interrupts this cascade by requiring that each information stream maintain its own confidence classification:

- Intelligence assessments must be labeled with ASE band and source count.

- Political statements in formal ACP channels must reference the ASE classification, not assert independent certainty.
- Media monitoring (via the ACP Narrative Commons) tracks narrative divergence from verified attribution levels.

Certainty cannot migrate between streams without an explicit translation step that preserves the original confidence classification. This does not prevent political actors from making public claims — it prevents those claims from being recognized as procedurally valid within the ACP system until the ASE threshold is met.

A.5 External Attribution Mediation Node (EAMN)

ICL requires a neutral, technically credible third party to aggregate evidence streams and publish probabilistic attribution assessments. This function is not a court or arbitration body — it issues non-binding confidence syntheses that participants may use to calibrate their own assessments.

Suitable EAMN hosts include existing institutions with technical verification capacity: the IAEA (for nuclear-adjacent incidents), UN-affiliated monitoring bodies, or designated neutral-state verification units. The EAMN does not determine guilt; it publishes evidence-weighted confidence bands.

A.6 Integration with ACP v0.5

ICL operates as a pre-condition layer before ACP enters Emergency Mode. The integration point is explicit:

1. A high-impact incident occurs.
2. ICL activates automatically; LDW begins.
3. ASE classifies all attribution claims entering ACP channels.
4. ACP remains in Containment-State Coordination Mode until one of the following: (a) LDW expires without escalation, or (b) ASE reaches Verified classification.
5. Only then does ACP transition to Emergency Mode with full escalation legitimacy.

This sequence prevents ACP's Emergency Mode from being triggered by politically manufactured certainty rather than verified facts.

A.7 Failure Modes and Design Honesty

ICL does not guarantee de-escalation. It is designed to convert irreversible cascade risk into staged decision pressure. Four failure modes are acknowledged:

Failure Mode	Description
Sovereignty override	Actors declare de facto certainty and act outside the LDW; ICL degrades to advisory status
Parallel reality formation	Competing intelligence ecosystems bypass formal ASE classification
Political time compression	Domestic pressure shortens the effective LDW below its formal duration
Narrative weaponization	Actors selectively cite confidence bands to justify pre-existing intent

In all failure modes, ICL still generates a documented record of the attribution trajectory — which feeds the Contribution Registry and CIL (Addendum B) for post-crisis accountability purposes.

Addendum B — Pre-Action Commitment Layer (PACL)

Status	Draft v0.1 — Action Threshold Governance
Supplements	ACP v0.5
Purpose	Bind escalation behavior to pre-agreed probabilistic attribution thresholds, and provide a rigorous threat model for how those thresholds fail under adversarial conditions

B.1 The Problem

ACP v0.5 tracks what actors do after commitments are made (Contribution Registry) and ICL governs how certainty forms before action. Neither governs the action threshold itself: the moment at which an actor decides that enough certainty exists to justify escalation.

PACL addresses this by asking actors to pre-declare, before any crisis occurs, what level of attribution confidence they require before taking each class of action. These declarations are not legally binding — they are publicly registered expectations that create reputational and procedural friction when violated.

B.2 Probabilistic Action Thresholding (PAT)

Each ACP participant registers a Pre-Action Commitment Contract (PCC) specifying the minimum ASE confidence band required to authorize each class of response:

Action Class	Default Minimum Threshold	Notes
Formal diplomatic protest	Medium ($\geq 40\%$)	Lowest threshold; minimal irreversibility
Economic or sanctions measures	High ($\geq 70\%$)	Significant but reversible impact
Military positioning or readiness change	High ($\geq 70\%$)	Escalatory signal; requires strong evidence
Targeted kinetic response	Verified ($\geq 85\%$)	Near-irreversible; highest threshold
Large-scale escalation	Verified ($\geq 95\%$)	Reserved for unambiguous, multi-source confirmed events

Participants may set thresholds above these defaults but not below them. Thresholds are registered publicly on the ACP Contribution Registry and reviewed annually. Divergence between declared thresholds and observed behavior is tracked (see B.4).

B.3 Sovereign Override Clause

Each participant retains an unconditional right to invoke a National Emergency Override, authorizing action below declared thresholds. The override does not require ACP approval. It does require:

- Public declaration within 6 hours of the triggering action.
- A written justification submitted to the ACP Secretariat within 72 hours.
- An automatic review by the Peer Mediation Panel within 14 days.

Override frequency and justification quality are tracked in the Contribution Registry. The override is not a violation — it is a named exception with documented consequences for trust weighting in future coordination cycles.

B.4 Adversarial Threat Model

PACL was stress-tested against six attack vectors. The threat model is published here as part of the protocol specification, not as an appendix, because practitioners need to understand the system's limits before deciding whether to rely on it.

Attack Vector	Severity	Description and Mitigation Limit
Threshold inflation at design time	Medium	Actors declare conservative thresholds with no intent to honor them. Mitigation: behavioral consistency tracking (B.4); but violations are only visible after action, creating a first-incident lag.
Crisis reinterpretation (semantic drift)	High	Actors redefine 'verified' or 'strong evidence' during a crisis to fit available intelligence. Mitigation: ASE definitions are locked at registration; divergence is flagged but not preventable.
Attribution laundering	Very High	Actors manufacture or amplify intelligence to push attribution into higher confidence bands. Mitigation: EAMN independence; but sufficiently sophisticated deception can bypass external verification.
Selective threshold revelation	High	Actors disclose only favorable commitments in diplomatic contexts while maintaining aggressive internal doctrine. Mitigation: full PCC publication is required; but internal doctrine remains unverifiable.
Override normalization	Medium-High	Repeated use of the sovereign override erodes its reputational cost over time. Mitigation: override frequency scoring with compounding trust discount; no hard limit.
Proxy attribution chaos	Critical	Proxy actors conduct operations specifically to keep attribution below threshold, enabling principals to act on 'insufficient' evidence without PACL legitimacy cost. No mitigation currently exists for this vector.

Note on the proxy vector: this is the most significant unresolved vulnerability in PACL. Modern conflict environments are partly structured to exploit attribution ambiguity. PACL slows the

legitimization of escalation but cannot close the gap between proxy behavior and principal accountability.

B.5 Integration with ICL and ACP v0.5

PACL and ICL operate in sequence:

1. ICL produces an ASE confidence classification for a given incident.
2. PACL maps that classification against the actor’s registered PCC thresholds.
3. If the ASE band falls below the relevant threshold, escalation is classified as procedurally unauthorized within ACP channels.
4. If the actor acts anyway, the override clause activates and the divergence is logged.

This sequence means that both layers must be active and interoperable for either to function effectively. A participant who registers with PACL but opts out of ICL’s ASE classification loses the shared reference point that PACL depends on.

Addendum C — System Regeneration Protocol (SRP)

Status	Draft v0.1 — Post-Fracture Re-Coherence Framework
Supplements	ACP v0.5
Purpose	Restore minimum viable coordination capacity after systemic fragmentation — without requiring agreement on what caused the fracture or who was responsible

C.1 The Problem

ACP v0.5, ICL, and PACL are all designed to prevent or contain breakdown. SRP addresses the condition after prevention has failed: when coordination architecture has fragmented, actors have exited shared frameworks, and the system needs to rebuild without pretending nothing happened.

This is not a hypothetical scenario. Coordination systems fragment regularly — after major attribution disputes, after sovereignty override cycles, after conflicts where one or more

parties declare that existing frameworks are illegitimate. SRP is the protocol for what comes next.

C.2 Design Principle: Interoperability, Not Unity

SRP does not attempt to restore a unified coordination architecture. It does not require actors to agree on what happened, who was responsible, or whose interpretation of events was correct. These questions may never be resolved.

SRP instead targets a narrower and more achievable goal:



GOAL *Rebuild sufficient shared reference points to enable partial coordination across incompatible governance realities — without requiring either side to concede their account of the fracture.*

This distinction matters operationally. Most post-conflict coordination attempts fail because they are structured as implicit reconciliation processes — requiring actors to behave as if shared frameworks are legitimate before those frameworks have been rebuilt. SRP separates coordination restoration from narrative resolution.

C.3 Four-Phase Regeneration Sequence

PHASE 1 — FRACTURE ACKNOWLEDGEMENT

The system formally declares a Multi-Reality State (MRS): an explicit recognition that ACP coordination has partially or fully broken down, that attribution systems are inconsistent across actors, and that existing frameworks are not currently functioning as designed.

This declaration is published on the ACP dashboard. It is not an admission of failure — it is a diagnostic statement that enables honest assessment of what can and cannot be rebuilt.

PHASE 2 — MINIMUM SHARED MODEL (MSM)

The MSM is the only element of SRP that is genuinely novel relative to existing ACP infrastructure. It defines a thin layer of shared reality constructed exclusively from:

- Observable physical events (documented incidents with time, location, and impact data),
- Verified infrastructural constraints (supply routes, border crossings, communication channels that exist or do not exist),

- Non-contested logistical facts (populations in specific locations, humanitarian needs with independent verification).

The MSM deliberately excludes: attribution, intent, responsibility, and interpretive framing. It contains only what all parties can observe independently. This is not a compromise position — it is an intentionally impoverished shared reality designed to be acceptable even to parties with incompatible accounts of what happened.

The MSM is the foundation on which everything else in SRP is built. Without it, no subsequent coordination is possible. With it, even deeply adversarial actors can coordinate on specific, bounded, observable tasks.

PHASE 3 — INTEROPERABILITY BRIDGING

Once an MSM exists, SRP introduces translation interfaces between the governance logics of actors operating under different frameworks. These are not merger mechanisms — they allow systems that disagree to coordinate partially without either adopting the other's framework.

In practice this means: Minimal Coordination Contracts (MCCs) — ultra-light agreements scoped to specific observable tasks (humanitarian corridor operation, communication channel reactivation, incident deconfliction in a defined area). MCCs are designed to be acceptable to actors who reject all broader coordination frameworks, because they make no claims about legitimacy, responsibility, or precedent.

PHASE 4 — STABILITY GRADIENT RECONSTRUCTION

Coordination capacity is rebuilt incrementally along a non-linear ladder. Regression is explicitly permitted — the ladder does not assume forward progress:

Stability Level	Description
Signal stability	Shared data streams re-established; actors receive the same observable facts
Operational stability	Logistics coordination resumes in defined, bounded domains
Crisis stability	Partial joint response capability for acute emergencies
Strategic stability	Limited trust-based coordination emerges; formal ACP frameworks become viable again

C.4 Residual Trust Ledger (RTL)

SRP introduces a Residual Trust Ledger — a record of reliability under crisis conditions that persists through fragmentation cycles. The RTL tracks:

- Past commitment fulfillment during crisis coordination (from the ACP Contribution Registry),
- Override behavior history (from PACL records),
- MCC compliance during SRP phases.

The RTL does not produce a trust score or ranking. It produces a probabilistic coordination weighting: an assessment of how reliably a given actor has honored bounded, specific commitments in the past. This weighting influences — but does not determine — how MCCs are structured in subsequent SRP cycles.

The RTL is designed to be compatible with the anticipated ITL (Informal Trust Ledger) infrastructure referenced in ACP v0.5, Section 5.2.2.

C.5 What SRP Cannot Do

SRP is an honest protocol. It explicitly acknowledges three conditions under which it provides no recovery path:

Condition	Why SRP Cannot Address It
Permanent fragmentation	Some actors may permanently exit shared coordination frameworks. SRP cannot compel re-entry; it can only reduce the cost of voluntary re-engagement.
Narrative divergence persistence	Actors may share observable facts while maintaining permanently incompatible interpretations. SRP does not require narrative convergence and cannot produce it.
Asymmetric recovery	Actors with greater institutional capacity may rebuild coordination capacity faster, creating imbalances that SRP's non-coercive structure cannot correct.

C.6 Integration with ACP v0.5

SRP activates when ACP formally declares a Multi-Reality State. It does not replace ACP — it creates the conditions under which ACP frameworks can become viable again. The relationship is sequential: SRP precedes ACP re-engagement, not the reverse.

The MSM data infrastructure is designed to be compatible with the ACP Situational Awareness Layer (SAL), allowing MSM-verified facts to migrate directly into SAL once strategic stability is re-established.

Full Stack Reference

These addenda complete the pre-coordination layer of the ACP governance stack. The full architecture, in operational sequence:

Layer	Component	Function
Pre-coordination	ICL (Addendum A)	Slows narrative certainty formation; prevents premature escalation trigger
Pre-coordination	PACL (Addendum B)	Governs action thresholds before crisis; documents divergence
Active coordination	ACP v0.5	Allocates resources, manages participation tiers, tracks contributions
Post-fracture	SRP (Addendum C)	Rebuilds minimum coordination capacity after systemic breakdown

The system does not guarantee de-escalation. It converts unstructured escalation risk into staged, observable, partially reversible decision sequences — and provides a recovery path when those sequences fail.