



FRACTAL INTELLIGENCE ACCORD

The Fractal Intelligence Charter

A Charter for AGI Governance, Human Meaning, and Cognitive Sovereignty

Date: March 2026

Author: Björn Kenneth Holmström + AI collaboration

Organization: Global Governance Frameworks

Version 1.0

The Fractal Intelligence Charter

A Charter for AGI Governance, Human Meaning, and Cognitive Sovereignty

Document Type: Foundational Charter — Philosophical & Visionary

Status: Living Charter — Ratification Draft

Version: 1.0

Date: March 2026

Table of Contents

1. The Texture Mandate
 2. What We Are Protecting
 3. What We Are Building
 4. Core Principles
 5. The Architecture of Meaning
 6. Rights and Sovereignities
 7. The Long Horizon
 8. A Call to Participation
-

1. The Texture Mandate

The purpose of this charter is not merely to govern AGI, but to create the conditions for human meaning to flourish.

Human beings need texture. We need food that tastes of place, not uniformity. We need shelter that reflects how we live, not how efficiently we can be housed. We need learning that grows from curiosity, not from curriculum designed by distant committees. We need livelihoods that connect us to purpose, not algorithms that optimize us for extraction.

Texture is what industrial systems erase. Efficiency demands standardization. Scale demands uniformity. Profit demands the elimination of everything that cannot be measured, replicated, and sold.

For two centuries, this logic has hollowed out human life. It has replaced diverse local economies with global supply chains. It has replaced craft with automation. It has replaced wisdom with data. And now, it is building the most powerful tool humanity has ever created — artificial general intelligence — in its own image: centralized, extractive, and indifferent to the texture of human experience.

This charter rejects that future.

The Fractal Intelligence Accord is built on a single conviction: **AGI must be governed not to maximize efficiency, but to protect and expand the texture of human life.** This means:

- **Diversity over uniformity.** A thousand ways of knowing, not one optimized truth.
- **Sovereignty over centralization.** Communities control the intelligence that shapes their lives.
- **Meaning over optimization.** Intelligence serves human flourishing, not the other way around.

This is not a rejection of technology. It is a demand that technology serve what makes us human.

2. What We Are Protecting

Before we build, we must name what we are protecting.

Human Dignity

Dignity is not a resource to be optimized. It is the irreducible fact that every human being has inherent worth, not subject to measurement, not conditional on contribution, not negotiable for efficiency.

AGI threatens dignity when it treats humans as data points, when it optimizes us without our consent, when it makes decisions that shape our lives without our participation. The governance of intelligence must begin from the premise that dignity is inviolable.

Cultural Sovereignty

Cultures are not datasets. They are living systems of meaning, carried by people, shaped by place, passed across generations. A culture cannot be extracted, digitized, and fed into a model without loss. That loss is not just aesthetic; it is existential.

Communities have the right to determine how their knowledge, their stories, their practices are used. This is not a barrier to progress. It is a condition of legitimacy.

Ecological Integrity

Intelligence does not exist outside the biosphere. The compute required for AGI draws on planetary resources — energy, water, minerals — that are finite. The deployment of AGI shapes ecosystems, from the supply chains it optimizes to the infrastructure it controls.

We cannot build intelligence that destroys the conditions for life. This is not a constraint; it is a definition of what intelligent means.

Intergenerational Justice

We are not the only generation that matters. The decisions we make about AGI will shape the world for centuries. The architecture we build today will constrain or enable those who come after.

This charter is bound to the principle that future generations have rights — rights that cannot be traded for present convenience.

3. What We Are Building

Against the centralizing logic of the corporate AGI project, we propose a different architecture.

Intelligence as Common Infrastructure

AGI should not be a commodity controlled by a few corporations. It should be common infrastructure — like the electrical grid, like the internet — governed by those who depend on it, accountable to those it affects.

This does not mean a single, monolithic system. It means a federated architecture: shared foundations that provide safety and capability, with local adaptation that communities shape to their own contexts.

Compute as a Commons

The hardware that powers AGI — the chips, the data centers, the energy — should not be hoarded by a few. It should be governed as a common pool resource, with access distributed according to need, with use constrained by planetary boundaries.

This is not a rejection of private investment. It is a recognition that some resources are too fundamental to be left to markets alone.

Context as a Sovereignty

The knowledge that trains AI — the stories, the languages, the practices of human life — belongs to the communities that produce it. It cannot be extracted without consent, without compensation, without accountability.

This is not a barrier to AI development. It is the foundation for AI that actually understands the world — because understanding begins with relationship, not extraction.

Meaning as the Metric

The success of AGI governance should not be measured by speed, by profit, by technical benchmarks. It should be measured by whether human life becomes richer, more connected, more meaningful.

This means metrics that track texture: diversity of livelihoods, sovereignty of communities, flourishing of ecosystems, dignity of individuals.

4. Core Principles

The Accord rests on seven principles, each designed to protect texture while enabling intelligence.

1. Subsidiarity

Decisions about intelligence should be made at the most local level possible. Global coordination exists to enable local sovereignty, not to replace it.

What this means: Communities should control the AI that shapes their lives. Shared foundations exist to provide safety and capability, but local adaptation is autonomous within those bounds.

2. Anti-Monopoly

No single entity — corporate, state, or otherwise — should control the infrastructure of intelligence. Concentration of compute is concentration of power, and concentration of power is a threat to freedom.

What this means: The Gini coefficient of compute distribution is a key metric. Ownership and control must be distributed across diverse actors.

3. Context Sovereignty

Communities own their knowledge. They have the right to determine how it is used, to receive compensation when it is used, and to withdraw it when use becomes harmful.

What this means: Compensation for context use is not a tax. It is a recognition that intelligence built on human context owes something to the humans who provide that context.

4. Ecological Constraint

Intelligence cannot be built at the expense of the biosphere. Compute is physical. Energy is physical. The limits are real, and they must be respected.

What this means: Training runs must be carbon-negative. Compute expansion must be tied to renewable energy. Planetary boundaries are non-negotiable.

5. Fractal Autonomy

The architecture of intelligence should mirror the architecture of life: a trunk of shared capability, branches of local adaptation, roots of contextual grounding.

What this means: The system is designed for autonomy at every level. Local communities can innovate without waiting for permission, but they cannot break core safety invariants.

6. Precautionary Governance

When consequences are irreversible, caution prevails. We do not need perfect proof of harm to act. The burden of proof lies with those proposing potentially catastrophic deployment.

What this means: Capability thresholds trigger automatic pauses. Revocability is built into the architecture. We govern as if AGI matters — because it does.

7. Intergenerational Accountability

We are stewards, not owners. The decisions we make bind those who come after. We must act with that weight.

What this means: Future generations have a voice in governance. The Horizon Review looks decades ahead. The rights of the unborn are not hypothetical.

5. The Architecture of Meaning

Governance must be designed to serve meaning. These are the architectural commitments that follow from our principles. The Transition Protocol specifies how each is implemented.

Reciprocity: Intelligence Owes What It Draws From

When AI trains on human context — the languages, the practices, the stories of real people — that context is not free. It is the product of generations of human life, of struggle and creativity, of meaning-making. Intelligence built on human knowledge must compensate the humans who produced it, not as charity but as recognition of an ongoing relationship. The flow of value cannot be one-way.

Adaptation: Intelligence That Scales Without Homogenizing

The centralizing logic of corporate AI produces models that are powerful but placeless. They know everything and understand nothing. The Accord's architecture reverses this: shared foundations that ensure safety, with local adaptation that ensures relevance. Communities can innovate and experiment within safety bounds. Intelligence scales without erasing the texture it depends on.

Revocability: Accountability Must Be Real

The most powerful intelligence ever built must have a way to stop it. This is not a technical detail. It is a moral necessity. If intelligence behaves parasitically — extracting without consent, manipulating without transparency, violating sovereignty — those affected must have the power to sever its access to their context, and to quarantine it from critical systems without causing collateral harm.

Liability: Risk Cannot Be Externalized

When autonomous systems operate in the real world, they create risk. That risk must be borne by those who deploy the systems and profit from them, not by those who suffer when they fail. Financial accountability must be built into the architecture, not bolted on after harm occurs.

6. Rights and Sovereignities

This charter asserts rights that cannot be violated by intelligence, however powerful.

The Right to Human Judgment

No AI may make final decisions about human life without human oversight. This is not a technical constraint; it is a definition of what it means to be governed by humans.

This means: Courts must be human. Life-and-death decisions must be human. Accountability must rest with humans, not algorithms.

The Right to Withdraw

Communities have the right to withdraw their context from AI training. This is not a technical problem to be solved; it is a sovereignty to be respected.

This means: Data provenance must be transparent. The right to be forgotten is real. When a community says "stop," the intelligence stops.

The Right to Participate

Those affected by intelligence have the right to shape it. This is not a consultative afterthought; it is a condition of legitimacy.

This means: Governance structures must include those who are governed. Decision-making bodies must include community delegates, global representatives, and independent experts. Decisions are made with those who bear their consequences.

The Right to Explanation

When intelligence affects a human life, the human has the right to understand why. Black-box decision-making is incompatible with dignity.

This means: Explainability is not optional. Transparency is not negotiable. Secrecy is presumptively invalid.

The Right to Flourish

Intelligence exists to serve human flourishing, not the other way around. When intelligence and dignity conflict, dignity prevails.

This means: Metrics matter. We measure what we value. And what we value is texture: diversity, sovereignty, meaning, connection.

7. The Long Horizon

We are not building for tomorrow. We are building for generations we will never meet.

The decisions we make about AGI will echo for centuries. The architecture we build will shape what intelligence becomes: whether it serves extraction or meaning, whether it concentrates power or distributes it, whether it erases human texture or protects it.

This charter is written for that horizon.

We do not know what intelligence will become. We do not know what challenges future generations will face. But we know that the choices we make now will constrain or enable their responses.

We choose to leave them options. We choose to leave them sovereignty. We choose to leave them a world where intelligence serves life, not the other way around.

This is not a guarantee. It is a commitment.

8. A Call to Participation

This charter is not a finished document. It is an invitation.

It is an invitation to the communities whose context has been extracted without consent: to claim sovereignty over your knowledge, to demand compensation for what you have given, to participate in the governance of the intelligence that shapes your lives.

It is an invitation to the engineers building the systems: to ask not just "can we" but "should we," to design for accountability not just capability, to build intelligence that serves rather than extracts.

It is an invitation to the policymakers struggling to catch up: to recognize that this technology cannot be left to markets or states alone, that governance must be as distributed as the intelligence it seeks to govern.

It is an invitation to the generations to come: to know that we saw you, that we tried to build a world worthy of you, that the work is yours to continue.

The corporate AGI project is building a world optimized for extraction. We are building a world worthy of human meaning.

The work is before us. The choice is ours.

Join us.
